

230962258_Aadiv_49_ML-3

August 5, 2025

Machine Learning Lab #3

DATA PREPROCESSING AND REGRESSION

[]:

[18]: *#Generic Imports*

```
import numpy as np
import pandas as pd
import matplotlib
from matplotlib import pyplot as plt
from scipy.sparse import csr_matrix
import seaborn as sns
```

[19]:

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix
```

Questions

1. Consider the hepatitis/ pima-indians-diabetes csv file, perform the following data pre-processing.

1. Load data in Pandas.

[20]:

```
df = pd.read_csv("/home/cv1-5aiml-a2/Documents/ML_230962258_AIML_A/Datasets/
↳diabetes_csv.csv")
```

2. Drop columns that aren't useful.

[21]: *# Dropped below:*

```
df = df.drop("Pregnancies",axis=1)
df = df.drop("SkinThickness", axis = 1)
df.reset_index(drop=True, inplace=True)

df.head()
```

[21]:

	Glucose	BloodPressure	Insulin	BMI	DiabetesPedigreeFunction	Age	\
0	148	72	0	33.6	0.627	50	
1	85	66	0	26.6	0.351	31	
2	183	64	0	23.3	0.672	32	

3	89	66	94	28.1	0.167	21
4	137	40	168	43.1	2.288	33

	Outcome
0	1
1	0
2	1
3	0
4	1

3. Drop rows with missing values.

```
[22]: df.dropna(axis = 0)
```

```
[22]:      Glucose  BloodPressure  Insulin  BMI  DiabetesPedigreeFunction  Age \
0         148             72         0  33.6              0.627      50
1          85             66         0  26.6              0.351      31
2         183             64         0  23.3              0.672      32
3          89             66        94  28.1              0.167      21
4         137             40       168  43.1              2.288      33
..      ...
763       101             76       180  32.9              0.171      63
764       122             70         0  36.8              0.340      27
765       121             72       112  26.2              0.245      30
766       126             60         0  30.1              0.349      47
767        93             70         0  30.4              0.315      23
```

	Outcome
0	1
1	0
2	1
3	0
4	1
..	...
763	0
764	0
765	0
766	1
767	0

[768 rows x 7 columns]

4. Create dummy variables
5. Take care of missing data.
6. Convert the data frame to NumPy.

```
[23]: np_df = df.to_numpy()
```

7. Divide the data set into training data and test data.

```
[24]: X = df.drop('Outcome', axis=1)
      Y = df['Outcome']

      x_train = X[:80]
      y_train = Y[:80]
      x_test = X[80:]
      y_test = Y[80:]
```

2. a. Construct a CSV file with the following attributes:

Study time in hours of ML lab course (x)

Score out of 10 (y)

The dataset should contain 10 rows.

```
[37]: x = np.random.randint(0, 10, 10)
      y = np.random.randint(0, 10, 10)
      data = {'x': x, 'y': y}
      df = pd.DataFrame(data)
      df.to_csv('Q2a.csv', index=False)

      print("CSV file 'Q2a.csv' has been created.")
```

CSV file 'Q2a.csv' has been created.

b. Create a regression model and display the following:

Coefficients: B0 (intercept) and B1 (slope)

RMSE (Root Mean Square Error)

Predicted responses

```
[38]: df_read = pd.read_csv('Q2a.csv')
      print("\nFirst 5 rows of the dataset:")
      print(df_read.head())
      print("\nInformation about the dataset:")
      print(df_read.info())
```

First 5 rows of the dataset:

	x	y
0	7	1
1	5	7
2	9	5
3	2	9
4	5	6

Information about the dataset:

<class 'pandas.core.frame.DataFrame'>

```

RangeIndex: 10 entries, 0 to 9
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0    x         10 non-null    int64
 1    y         10 non-null    int64
dtypes: int64(2)
memory usage: 292.0 bytes
None

```

c. Create a scatter plot of the data points in red color and plot the graph of x vs. predicted y in blue color.

d. Implement the model using two methods:

Pedhazur formula (intuitive)

Calculus method (partial derivatives, refer to class notes)

```

[ ]: df = pd.read_csv('Q2a.csv')
x = df['x'].values
y = df['y'].values
n = len(x)

# Method 1: Pedhazur Formula (Intuitive/Analytical)
x_mean = np.mean(x)
y_mean = np.mean(y)

# slope
numerator_b1 = np.sum((x - x_mean) * (y - y_mean))
denominator_b1 = np.sum((x - x_mean)**2)
b1_pedhazur = numerator_b1 / denominator_b1
#intercept
b0_pedhazur = y_mean - b1_pedhazur * x_mean

# -----
# Method 2: Calculus Method (Partial Derivatives)
# -----

sum_x = np.sum(x)
sum_y = np.sum(y)
sum_x2 = np.sum(x**2)
sum_xy = np.sum(x * y)

A = np.array([[n, sum_x], [sum_x, sum_x2]])
B = np.array([sum_y, sum_xy])

try:
    solution = np.linalg.solve(A, B)

```

```

b0_calculus = solution[0]
b1_calculus = solution[1]
except np.linalg.LinAlgError:
    b0_calculus, b1_calculus = None, None
    print("Error: Matrix is singular. Cannot solve the system of equations.")

```

- e. Compare the coefficients obtained using both methods and compare them with the analytical solution.

```

[43]: print("Coefficients from Pedhazur Formula:")
      print(f"B0 (Intercept): {b0_pedhazur}")
      print(f"B1 (Slope): {b1_pedhazur}")
      print("\nCoefficients from Calculus Method:")
      if b0_calculus is not None:
          print(f"B0 (Intercept): {b0_calculus}")
          print(f"B1 (Slope): {b1_calculus}")

```

Coefficients from Pedhazur Formula:

B0 (Intercept): 5.447619047619048

B1 (Slope): -0.30952380952380953

Coefficients from Calculus Method:

B0 (Intercept): 5.447619047619046

B1 (Slope): -0.3095238095238093

- f. Test your model to predict the score obtained when the study time of a student is 10 hours.

```

[45]: study_time_new = 10
      predicted_score_new = b0_calculus + b1_calculus * study_time_new

      print(f"\nPrediction:")
      print(f"When the study time is {study_time_new} hours, the predicted score is_
      ↪{predicted_score_new:2f}.")

```

Prediction:

When the study time is 10 hours, the predicted score is 2.352381.
